

Kubernetes Resources Part 1

Summary

Up to this stage, we have explored how to package an application using Docker and how to bootstrap a cluster using k3s. In the next phase, we need to deploy the packaged application to a Kubernetes cluster.

Kubernetes provides a rich collection of resources that are used to deploy, configure, and manage an application. Some of the widely used resources are:


- Pods - the atomic element within a cluster to manage an application
- Deployments & ReplicaSets - oversees a set of pods for the same application
- Services & Ingress - ensures connectivity and reachability to pods
- Configmaps & Secrets - pass configuration to pods
- Namespaces - provides a logical separation between multiple applications and their resources
- Custom Resource Definition (CRD) - extends Kubernetes API to support custom resources

Application Deployment

Summary


A **pod** is the anatomic element within a cluster that provides the execution environment for an application. Pods are the smallest manageable units in a Kubernetes cluster. Every pod has a container within it, that executes an application from a Docker image (or any OCI-compliant image). There are use cases where 2-3 containers run within the same pod, however, it is highly recommended to keep the 1:1 ratio between your pods and containers.

All the pods are placed on the cluster nodes. A node can host multiple pods for different applications.



Pod architecture, showcasing a container running a Docker image

Deployments and ReplicaSets



Application management using a Deployment and ReplicaSet

To deploy an application to a Kubernetes cluster, a **Deployment** resource is necessary. A Deployment contains the specifications that describe the desired state of the application. Also, the Deployment resource manages pods by using a **ReplicaSet**. A ReplicaSet resource ensures that the desired amount of replicas for an application are up and running at all times.

To create a deployment, use the `kubectl create deployment` command, with the following syntax:

```
# create a Deployment resource
# NAME - required; set the name of the deployment
# IMAGE - required; specify the Docker image to be executed
```

```
# FLAGS - optional; provide extra configuration parameters for the resource  
# COMMAND and args - optional; instruct the container to run specific commands when it starts  
kubectl create deploy NAME --image=image [FLAGS] -- [COMMAND] [args]  
  
# Some of the widely used FLAGS are:  
-r, --replicas - set the number of replicas  
-n, --namespace - set the namespace to run  
--port - expose the container port
```

For example, to create a Deployment for the Go hello-world application, the following command can be used:

```
# create a go-helloworld Deployment in namespace `test`  
kubectl create deploy go-helloworld --image=pixelpotato/go-helloworld:v1.0.0 -n test
```

It is possible to create headless pods or pods that are not managed through a ReplicaSet and Deployment. Though it is not recommended to create standalone pods, these are useful when creating a testing pod.

To create a headless pod, the `kubectl run` command is handy, with the following syntax:

```
# create a headless pod  
# NAME - required; set the name of the pod  
# IMAGE - required; specify the Docker image to be executed  
# FLAGS - optional; provide extra configuration parameters for the resource  
# COMMAND and args - optional; instruct the container to run specific commands when it starts  
kubectl run NAME --image=image [FLAGS] -- [COMMAND] [args...]  
  
# Some of the widely used FLAGS are:  
--restart - set the restart policy. Options [Always, OnFailure, Never]  
--dry-run - dry run the command. Options [none, client, server]  
-it - open an interactive shell to the container
```

For example, to create a busybox pod, the following command can be used:

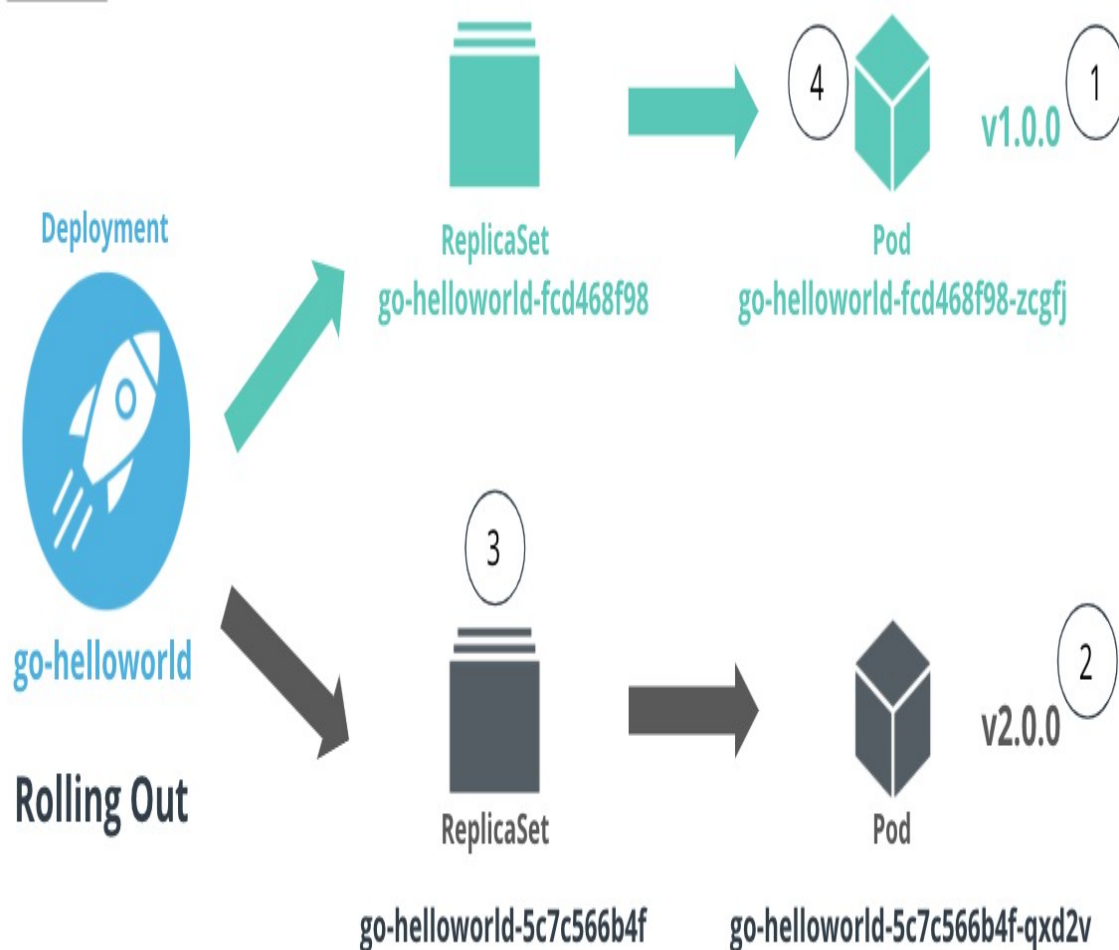
```
# example: create a busybox pod, with an interactive shell and a
restart policy set to Never
kubectl run -it busybox-test --image=busybox --restart=Never
```

Rolling Out Strategy

The Deployment resource comes with a very powerful rolling out strategy, which ensures that no downtime is encountered when a new version of the application is released. Currently, there are 2 rolling out strategies:

- RollingUpdate - updates the pods in a rolling out fashion (e.g. 1-by-1)
- Recreate - kills all existing pods before new ones are created

For example, in this case, we upgrade a Go hello-world application from version 1.0.0 to version 2.0.0:



Rolling update of an application, between different versions

Where:

1. The Go hello-world application is running version v1.0.0 in a pod managed by a ReplicaSet
2. The version of Go hello-world application is set to v2.0.0
3. A new ReplicaSet is created that controls a new pod with the application running in version v2.0.0
4. The traffic is directed to the pod running v2.0.0 and the pod with the old configuration (v1.0.0) is removed

Application Development Demo

Summary

This demo showcases how an application can be deployed, configured, and managed within a Kubernetes cluster using Deployment, ReplicaSet, and pod resources.

The instructor uses the Go hello-world application in version v1.0.0 and v2.0.0. The difference between these 2 versions is the exposed port by the application. Below are the code snippets for both application versions (you can also refer to the [go-helloworld](#) application from the course repository):

```
# Application version: v1.0.0
# port exposed: 6111
package main

import (
    "fmt"
    "net/http"
)

func helloWorld(w http.ResponseWriter, r *http.Request){
    fmt.Fprintf(w, "Hello World")
}

func main() {
    http.HandleFunc("/", helloWorld)
```

```
    http.ListenAndServe(":6111", nil)
}
# Application version: v2.0.0
# port exposed: 6112
package main

import (
    "fmt"
    "net/http"
)

func helloWorld(w http.ResponseWriter, r *http.Request){
    fmt.Fprintf(w, "Hello World")
}

func main() {
    http.HandleFunc("/", helloWorld)
    http.ListenAndServe(":6112", nil)
}
```

New terms

- **Pod** - smallest manageable unit within a cluster that provides the execution environment for an application
- **ReplicaSet** - a mechanism to ensure a number of pod replicas are up and running at all times
- **Deployment** - describe the desired state of the application to be deployed

Further reading

Explore the Kubernetes resources in more detail:

- [Kubernetes Pods](#)
- [Kubernetes Deployments](#)
- [Kubernetes ReplicaSets](#)
- [Kubernetes RollingOut Strategies](#)